

Distributed Matrix Multiplication

Venkata Paredla
201551035

Introduction

- Multiplication between two matrices requires huge computational power and also requires more memory, when the size of matrix is large.
- When size of matrices become larger, it becomes inefficient and time consumption task.

Problem Definition

Given a matrix A ($m \times n$) m rows and n columns, where each of its elements is denoted a_{ij} With $1 \leq i \leq m$ and $1 \leq j \leq n$, and a matrix B ($r \times n$) of r rows and n columns, Where each of its elements is denoted b_{ij} with $1 \leq i \leq r$, and $1 \leq j \leq n$, the matrix C resulting from the operation of multiplication of matrices A and B , $C = A \times B$, is such that each of its elements is denoted c_{ij} with $1 \leq i \leq m$.

Traditional Approach

- Calculated as $c_{ij} = \sum_{k=1}^r a_{ik} \times b_{kj}$
- The number of operations required to multiply A X B is $m \times n \times (2r - 1)$
- For simplicity, usually it is analyzed in terms of square matrices of order n. So that the quantity of basic operations between scalars is

$$2n^3 - n^2 = O(n^3)$$

Implementation of Parallel Algorithm:

Consider two square matrices A and B of size n that have to be multiplied:

1. Partition these matrices in square blocks p, where p is the number of processes available.
2. Create a matrix of size $p^{1/2} \times p^{1/2}$, so that each process can maintain a block of A matrix and a block of B matrix.
3. Each block is sent to each process, and the copied sub blocks are multiplied together and the results added to the partial results in the C sub-blocks.
4. The A sub-blocks are rolled one step to the left and the B sub-blocks are rolled one step upward.
5. Repeat steps 3 and 4 \sqrt{p} times.

Cannon's algorithm

```
for all (i=0 to s-1)      ... "skew" A
  Left-circular-shift row i of A by i,
    so that A(i,j) is overwritten by A(i, (j+i) mod s)
end for
for all (i=0 to s-1)      ... "skew" B
  Up-circular-shift column i of B by i,
    so that B(i,j) is overwritten by B( (i+j) mod s, j)

end for
for k=0 to s-1
  for all (i=0 to s-1, j=0 to s-1)
    C(i,j) = C(i,j) + A(i,j)*B(i,j)
    Left-circular-shift each row of A by 1,
      so that A(i,j) is overwritten by A(i, (j+1) mod s)
    Up-circular-shift each column of B by 1,
      so that B(i,j) is overwritten by B( (i+1) mod s, j)
  end for
end for
```

Performance

- Skewing A. $(s/2) * (\alpha + (n/s)^2 * \beta) = \sqrt{p} * \alpha / 2 + n^2 / (2 * \sqrt{p}) * \beta$
- Skewing B. The cost is the same as for skewing A.
- Shifting A (or B) left (or up) by 1. This costs $\alpha + n^2 / p * \beta$.
- Local accumulation of A times B into C. This costs $2 * (n/s)^3 = 2 * n^3 / p^{(3/2)}$.
- The total cost is therefore,

$$\text{Time} = 2 * n^3 / p + 3 * \sqrt{p} * \alpha + 3 * n^2 / \sqrt{p} * \beta$$

Assumptions

- For simplicity, worked with square matrices.
- Size of the matrix must be divisible by square root of available machines (for block division).

Implementation

- **Language:** Python
- **Libraries:** Numpy and Pyro4
- Client program takes following arguments
 - Available Machines
 - Dimension of matrices
- Server program takes following arguments
 - Port number

Scope

- Generalizing to matrix of any size.
- Synchronization when there are latency issues in the network.

Resources

1. <https://cse.buffalo.edu/faculty/miller/Courses/CSE633/Ortega-Fall-2012-CSE633.pdf>
2. <https://people.eecs.berkeley.edu/~demmel/cs267/lecture11/lecture11.html>

Thanks

- Source code is at https://github.com/raghu5910/distributed_matrix_multiplication